

```
In [15]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [16]: df = pd.read_csv("Air_Quality.csv")
df.head()
```

```
Out[16]:
```

	country	state	city	station	last_update	latitude	longitude
--	---------	-------	------	---------	-------------	----------	-----------

0	India	Arunachal_Pradesh	Naharlagun	Naharlagun, Naharlagun - APSPCB	06-04-2025 15:00:00	27.103358	93.67964
1	India	Assam	Byrnihat	Central Academy for SFS, Byrnihat - PCBA	06-04-2025 15:00:00	26.071318	91.87488
2	India	Assam	Byrnihat	Central Academy for SFS, Byrnihat - PCBA	06-04-2025 15:00:00	26.071318	91.87488
3	India	Assam	Guwahati	IITG, Guwahati - PCBA	06-04-2025 15:00:00	26.202864	91.70046
4	India	Assam	Guwahati	LGBI Airport, Guwahati - PCBA	06-04-2025 15:00:00	26.108870	91.58954



```
In [17]: # Display the first few rows
print(" ♦ First 5 Rows:")
display(df.head())
```

♦ First 5 Rows:

	country	state	city	station	last_update	latitude	longitude
0	India	Arunachal_Pradesh	Naharlagun	Naharlagun, Naharlagun - APSPCB	06-04-2025 15:00:00	27.103358	93.679645
1	India	Assam	Byrnihat	Central Academy for SFS, Byrnihat - PCBA	06-04-2025 15:00:00	26.071318	91.874880
2	India	Assam	Byrnihat	Central Academy for SFS, Byrnihat - PCBA	06-04-2025 15:00:00	26.071318	91.874880
3	India	Assam	Guwahati	IITG, Guwahati - PCBA	06-04-2025 15:00:00	26.202864	91.700464
4	India	Assam	Guwahati	LGBI Airport, Guwahati - PCBA	06-04-2025 15:00:00	26.108870	91.589544

In [18]: `# Check data structure`  
`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3187 entries, 0 to 3186
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   country         3187 non-null   object
1   state           3187 non-null   object
2   city            3187 non-null   object
3   station         3187 non-null   object
4   last_update     3187 non-null   object
5   latitude        3187 non-null   float64
6   longitude       3187 non-null   float64
7   pollutant_id    3187 non-null   object
8   pollutant_min   3046 non-null   float64
9   pollutant_max   3046 non-null   float64
10  pollutant_avg   3046 non-null   float64
dtypes: float64(5), object(6)
memory usage: 274.0+ KB
```

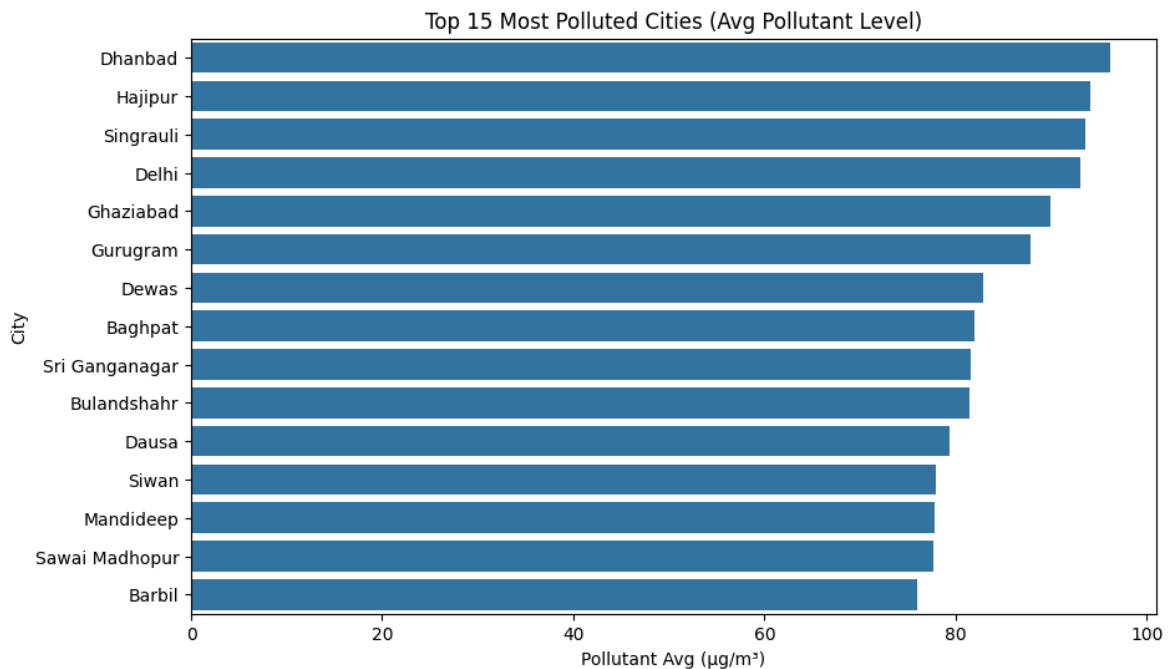
In [19]: `# Convert date column`  
`df['last_update'] = pd.to_datetime(df['last_update'], errors='coerce')`

In [20]: `# Check missing values`  
`df.isnull().sum()`

```
Out[20]: country      0
state      0
city      0
station    0
last_update 0
latitude   0
longitude  0
pollutant_id 0
pollutant_min 141
pollutant_max 141
pollutant_avg 141
dtype: int64
```

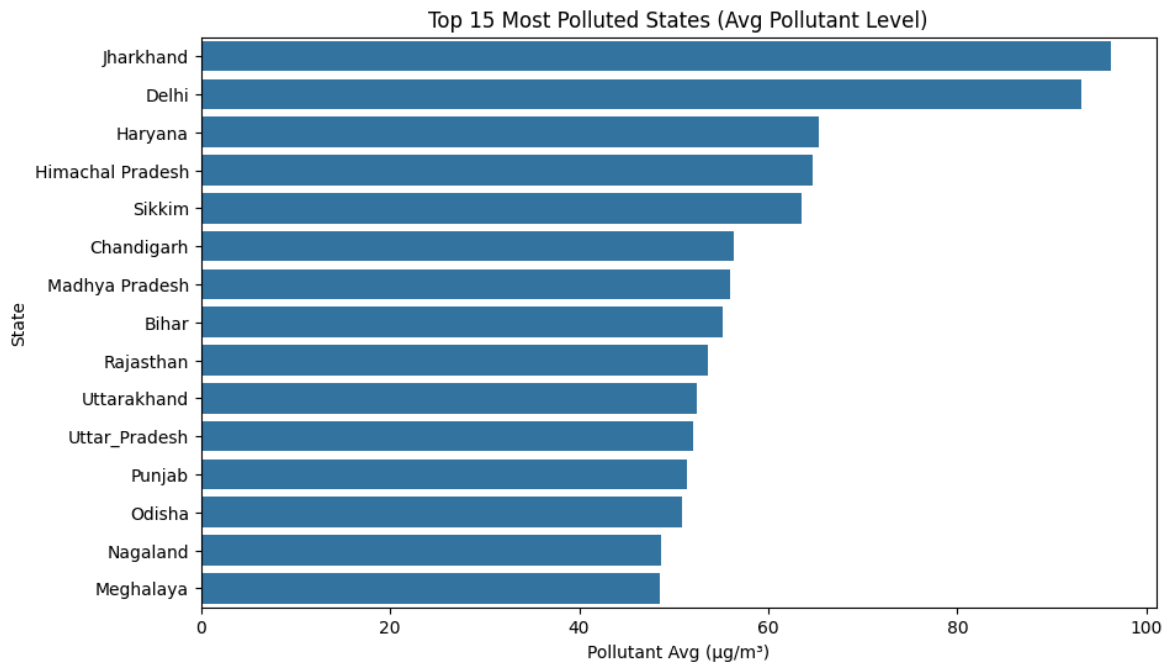
```
In [21]: # Average pollution per city
city_avg = df.groupby('city')['pollutant_avg'].mean().sort_values(ascending=False)

plt.figure(figsize=(10,6))
sns.barplot(x=city_avg.values, y=city_avg.index)
plt.title("Top 15 Most Polluted Cities (Avg Pollutant Level)")
plt.xlabel("Pollutant Avg ( $\mu\text{g}/\text{m}^3$ )")
plt.ylabel("City")
plt.show()
```



```
In [22]: # Average pollution per state
state_avg = df.groupby('state')['pollutant_avg'].mean().sort_values(ascending=False)

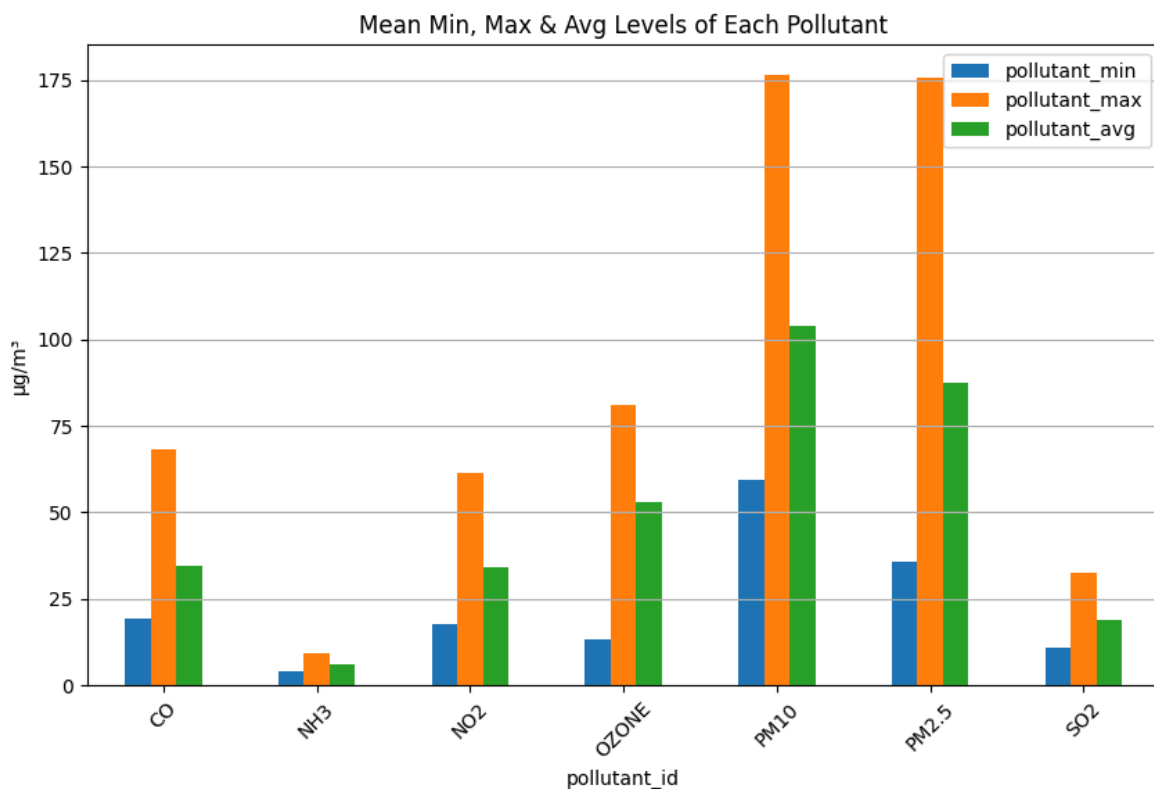
plt.figure(figsize=(10,6))
sns.barplot(x=state_avg.values, y=state_avg.index)
plt.title("Top 15 Most Polluted States (Avg Pollutant Level)")
plt.xlabel("Pollutant Avg ( $\mu\text{g}/\text{m}^3$ )")
plt.ylabel("State")
plt.show()
```



```
In [23]: # Group by pollutant type
pollutant_stats = df.groupby('pollutant_id')[['pollutant_min', 'pollutant_max',
print(pollutant_stats)

# Visualize
pollutant_stats.plot(kind='bar', figsize=(10,6))
plt.title("Mean Min, Max & Avg Levels of Each Pollutant")
plt.ylabel("µg/m³")
plt.xticks(rotation=45)
plt.grid(axis='y')
plt.show()
```

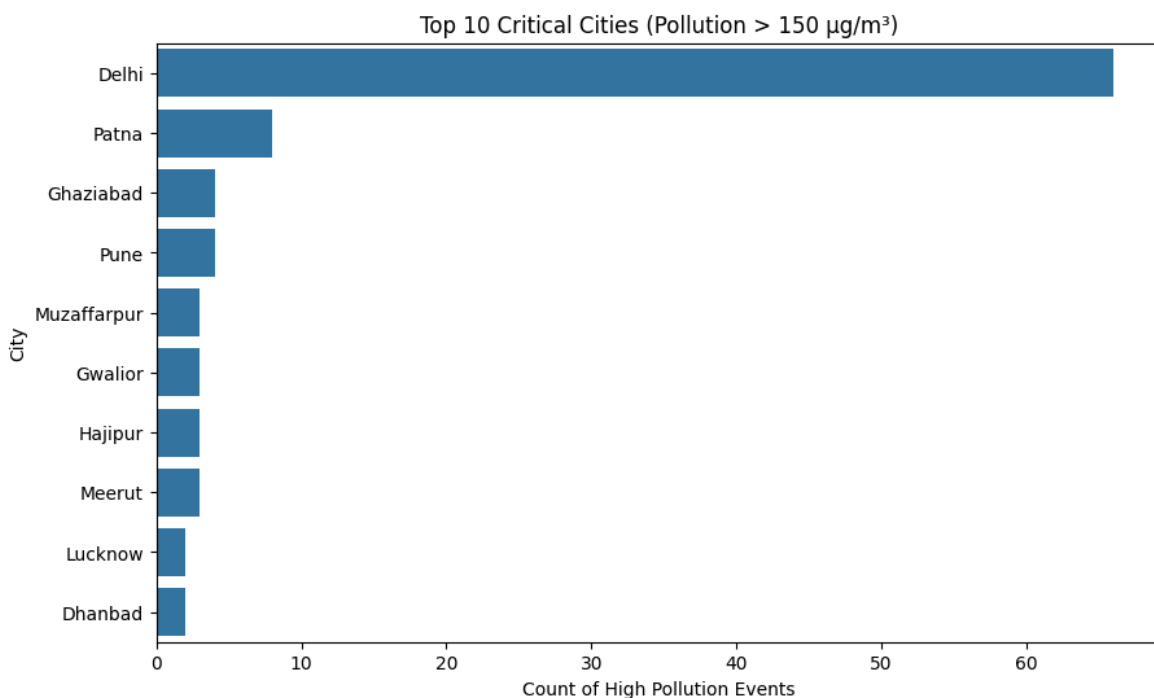
	pollutant_min	pollutant_max	pollutant_avg
pollutant_id			
CO	19.401330	68.334812	34.665188
NH3	4.137500	9.215000	6.020000
NO2	17.563063	61.461712	34.006757
OZONE	13.402299	80.928736	52.891954
PM10	59.358277	176.410431	103.795918
PM2.5	35.597315	175.648770	87.581655
SO2	11.016355	32.369159	18.738318



```
In [24]: # Define critical threshold
threshold = 150

# Cities with most high-pollution records
high_pollution = df[df['pollutant_avg'] > threshold]
critical_cities = high_pollution['city'].value_counts().head(10)

plt.figure(figsize=(10,6))
sns.barplot(x=critical_cities.values, y=critical_cities.index)
plt.title("Top 10 Critical Cities (Pollution > 150 µg/m³)")
plt.xlabel("Count of High Pollution Events")
plt.ylabel("City")
plt.show()
```

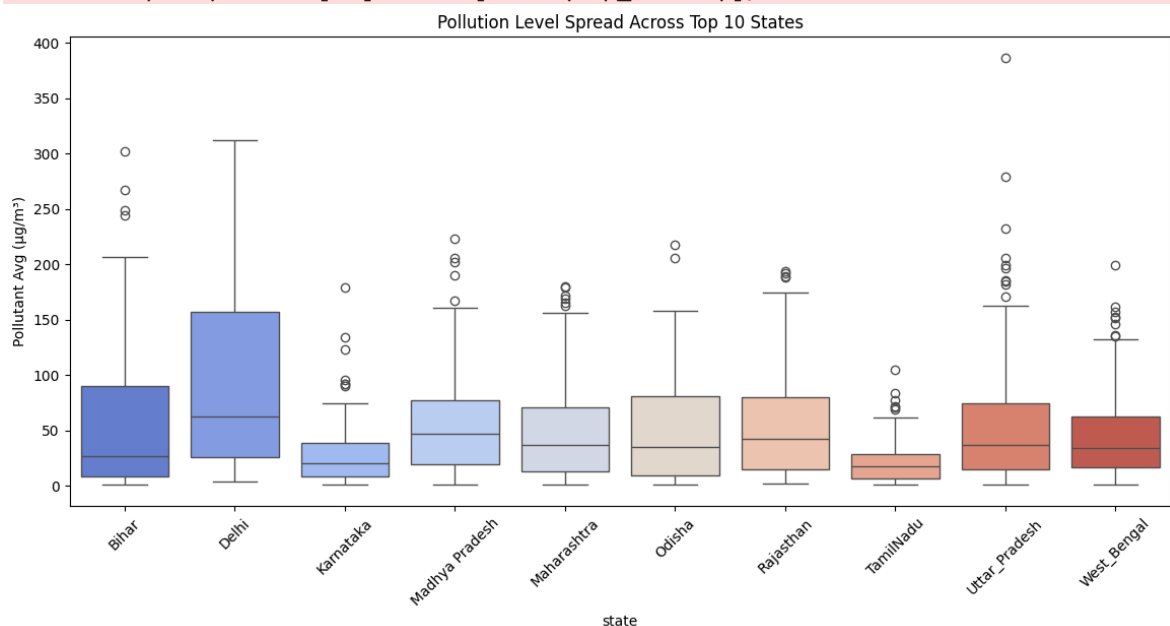


```
In [25]: # Boxplot to compare states
plt.figure(figsize=(14,6))
top_states = df['state'].value_counts().index[:10]
sns.boxplot(data=df[df['state'].isin(top_states)],
            x='state', y='pollutant_avg', palette='coolwarm')
plt.title("Pollution Level Spread Across Top 10 States")
plt.ylabel("Pollutant Avg ( $\mu\text{g}/\text{m}^3$ )")
plt.xticks(rotation=45)
plt.show()
```

C:\Users\SABITHA\AppData\Local\Temp\ipykernel\_10456\2056229655.py:4: FutureWarning:

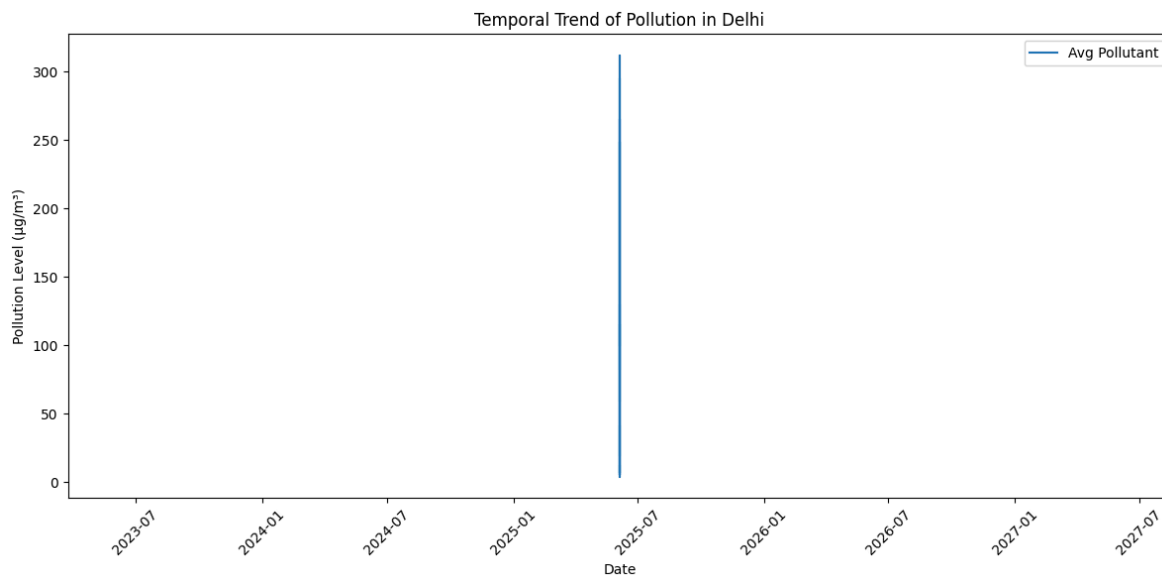
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df[df['state'].isin(top_states)],
```

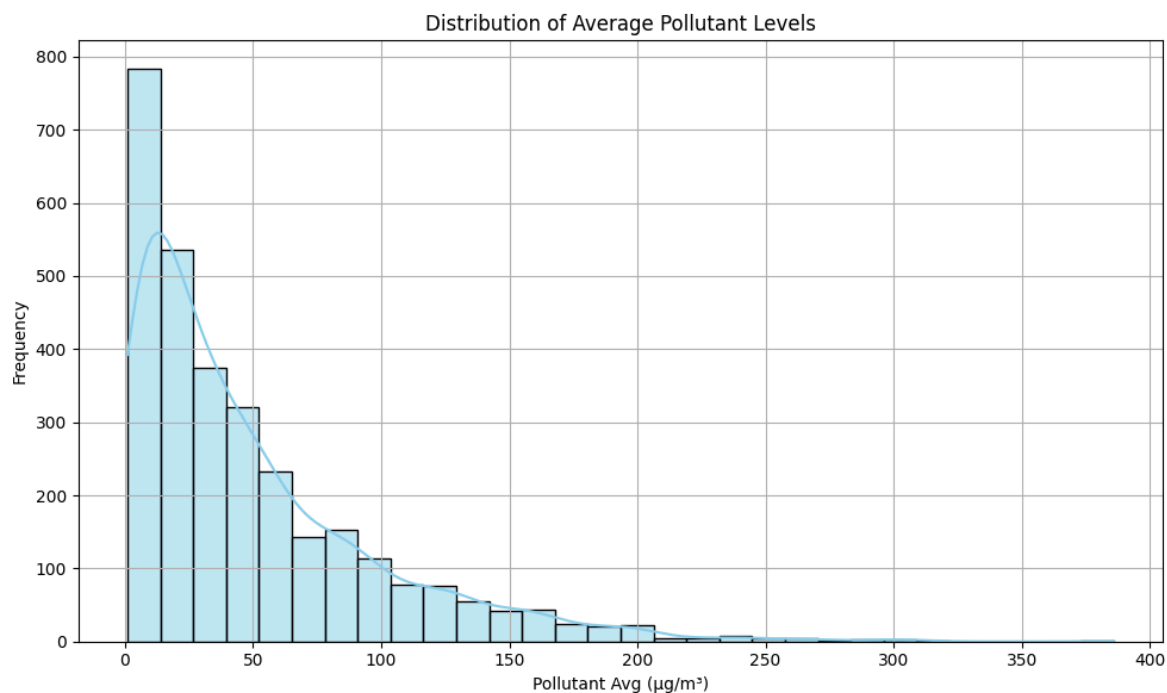


```
In [26]: # Choose a city (e.g., Delhi)
delhi = df[df['city'] == 'Delhi'].sort_values('last_update')


plt.figure(figsize=(12,6))
plt.plot(delhi['last_update'], delhi['pollutant_avg'], label='Avg Pollutant')
plt.title("Temporal Trend of Pollution in Delhi")
plt.xlabel("Date")
plt.ylabel("Pollution Level ( $\mu\text{g}/\text{m}^3$ )")
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()
```



```
In [27]: plt.figure(figsize=(10,6))
sns.histplot(df['pollutant_avg'], bins=30, kde=True, color='skyblue')
plt.title("Distribution of Average Pollutant Levels")
plt.xlabel("Pollutant Avg (µg/m³)")
plt.ylabel("Frequency")
plt.grid(True)
plt.tight_layout()
plt.show()
```




```
In [28]: plt.figure(figsize=(10,6))
sns.scatterplot(data=df, x='pollutant_avg', y='pollutant_max', hue='pollutant_id')
plt.title("Pollutant Avg vs Max Levels (by Pollutant Type)")
plt.xlabel("Average Pollutant Level (µg/m³)")
plt.ylabel("Maximum Pollutant Level (µg/m³)")
plt.legend(title='Pollutant', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```



```
In [29]: numeric_df = df[['pollutant_min', 'pollutant_max', 'pollutant_avg']]

# Compute correlation matrix
corr = numeric_df.corr()

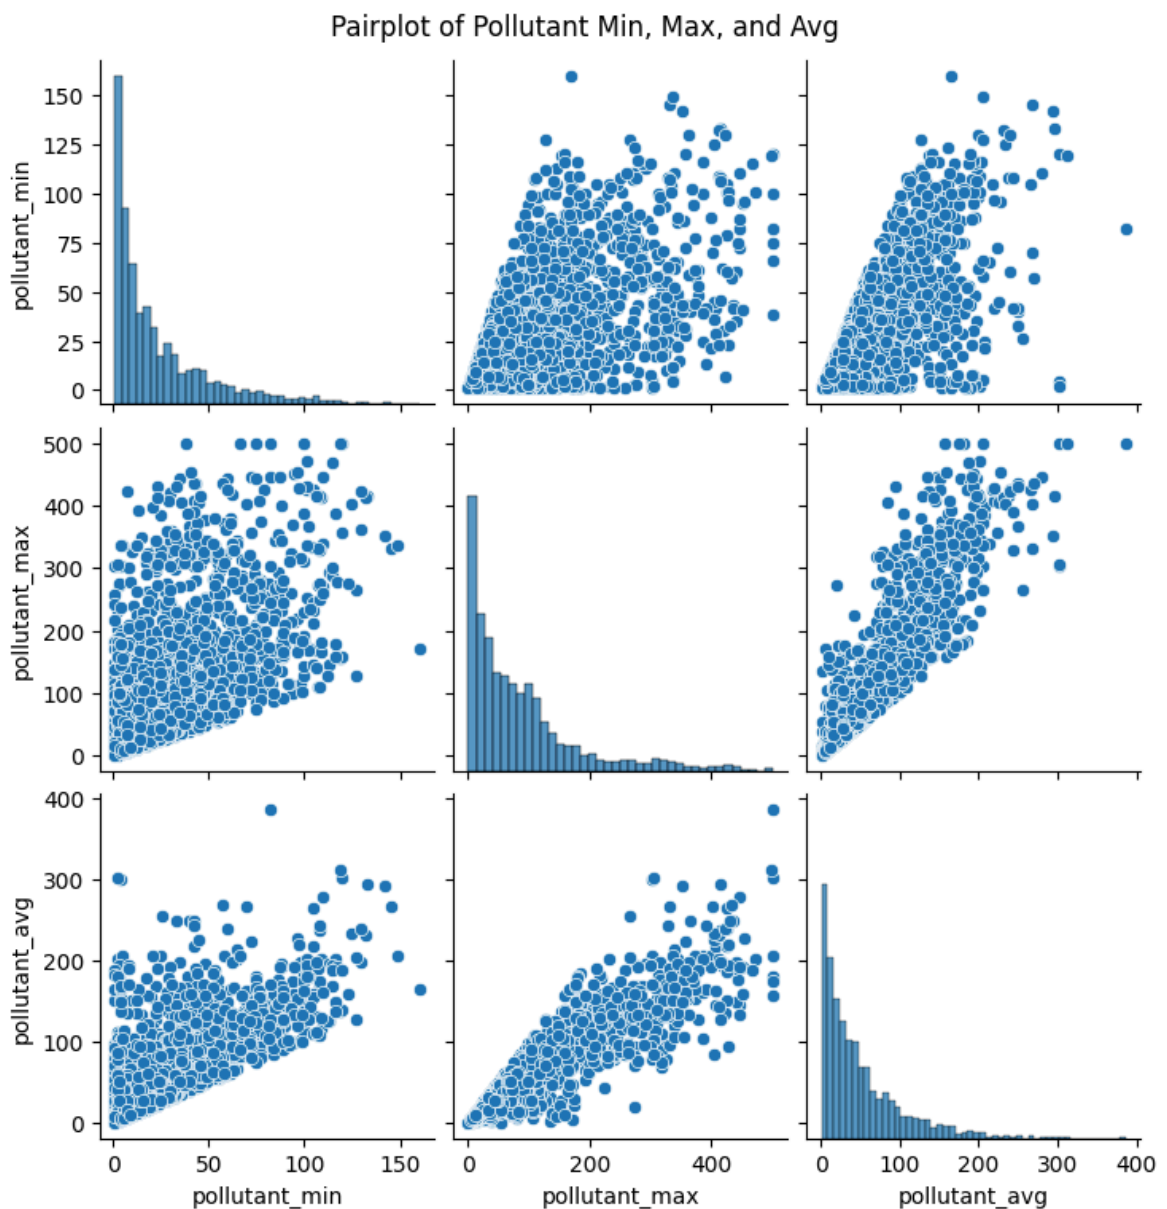
# Plot heatmap
plt.figure(figsize=(8,5))
sns.heatmap(corr, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title("Correlation Heatmap of Pollutant Levels")
plt.tight_layout()
plt.show()
```



```
In [30]: sns.pairplot(df[['pollutant_min', 'pollutant_max', 'pollutant_avg']])
plt.suptitle("Pairplot of Pollutant Min, Max, and Avg", y=1.02)
```



```
plt.show()
```

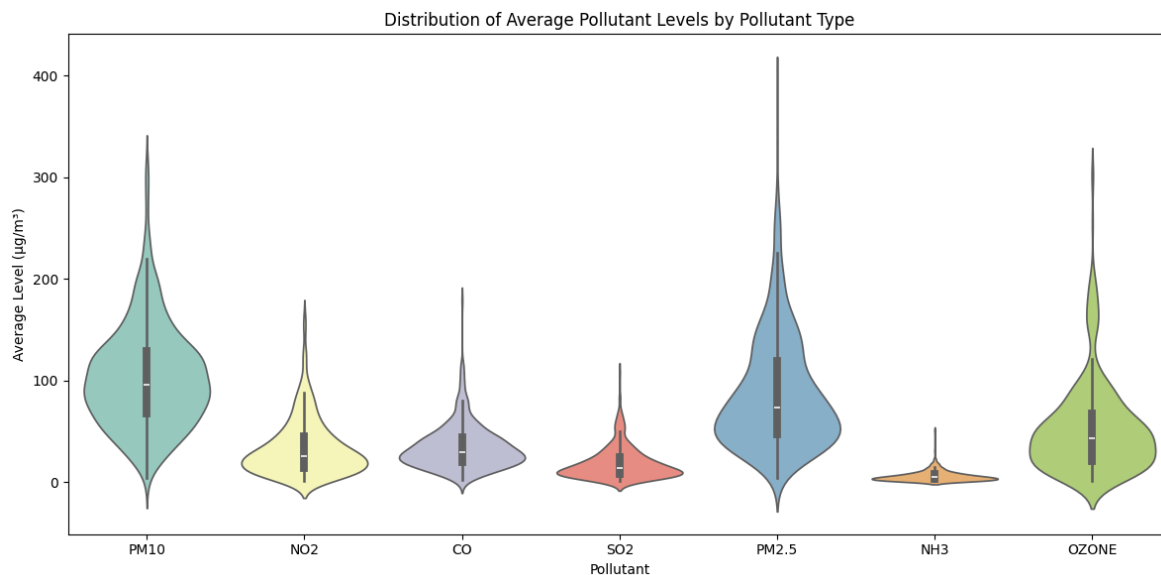


```
In [31]: plt.figure(figsize=(12,6))
sns.violinplot(data=df, x='pollutant_id', y='pollutant_avg', palette='Set3')
plt.title("Distribution of Average Pollutant Levels by Pollutant Type")
plt.xlabel("Pollutant")
plt.ylabel("Average Level ( $\mu\text{g}/\text{m}^3$ )")
plt.tight_layout()
plt.show()
```

C:\Users\SABITHA\AppData\Local\Temp\ipykernel\_10456\443648726.py:2: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.violinplot(data=df, x='pollutant_id', y='pollutant_avg', palette='Set3')
```




```
In [37]: pollutant_counts = df['pollutant_id'].value_counts()


plt.figure(figsize=(8,8))
plt.pie(pollutant_counts, labels=pollutant_counts.index, autopct='%1.1f%', startangle=90)
plt.title("Proportion of Pollutant Types Recorded")
plt.axis('equal') # Equal aspect ratio for a perfect circle
plt.show()
```



```
In [38]: sns.lmplot(data=df, x='pollutant_avg', y='pollutant_max', height=6, aspect=1.5,
plt.title("Linear Relationship: Avg vs Max Pollutant Levels")
plt.xlabel("Average Pollutant ( $\mu\text{g}/\text{m}^3$ )")
plt.ylabel("Max Pollutant ( $\mu\text{g}/\text{m}^3$ )")
plt.tight_layout()
plt.show()
```



```
In [39]: plt.figure(figsize=(10,6))
sns.kdeplot(data=df, x='pollutant_avg', fill=True, color='teal')
plt.title("KDE Plot of Average Pollutant Levels")
plt.xlabel("Pollutant Avg ( $\mu\text{g}/\text{m}^3$ )")
plt.ylabel("Density")
plt.grid(True)
plt.tight_layout()
plt.show()
```



In [ ]: